SHARKPIT

Master Product Brief & Build Blueprint

The Sparring Room for Your Next Big Moment

This document contains: Product Definition | Screen Specifications | Coaching Architecture | 1-Week Sprint Plan | Debate Strategy | Co-Founder Roles | Market Research

August 2026 — Internal Document

1. POSITIONING

SharkPit is not a presentation tool. It is not a rehearsal app. It is a **sparring room** for people who need to defend their ideas under live pressure.

You bring an idea — a thesis, a pitch, a proposal, a strategy. You enter a room with AI personas who challenge you, interrupt you, and probe your weak spots. After, you get coached: not scored, not graded, but **developed**. Over time, SharkPit learns your patterns and builds you toward the real room.

Tagline

> *"The sparring room for your next big moment."*

The Moment We Serve

Every person who creates anything — research, product, strategy, art — eventually stands in a room and must defend it. The grad student before a thesis committee. The founder before investors. The PM before leadership. The consultant before a client. Different people. Same moment. Same fear. Same skill gap.

Who It's For (And Not For)

For people who...	NOT for...
Have a specific idea they need to stand behind	Generic 'I want to be better at speaking'
Face a room that will ask hard questions	People giving one-way lectures
Are nervous because the stakes are real	People who just want slide design help
Need to practice recovery when interrupted	People practicing a script verbatim

2. NAVIGATION ARCHITECTURE

Four items. Clean. Purposeful. No more.

Nav Item	Purpose
Home	Your dashboard. Upcoming moment, today's focus, continue where you left off.
Practice	Start a sparring session. Upload material, configure the room, enter.
Coach	Your development layer. Active programs, skill map, pattern insights, recommended drills.
Progress	Session history, growth trends, your communication profile.

3. SCREEN DEFINITIONS

3.1 Home Screen

Feel: Like opening a premium training app. Calm, focused, purposeful.

- **Header:** 'Good evening, [Name].' Subtitle: *'Your [Thesis Defense] is in 14 days.'*
- **Primary CTA:** Big button — 'Enter the Sparring Room'
- **Continue Card:** Your last session or active program. One-tap resume.
- **Today's Focus (from Coach):** 'Last time you hedged when challenged on data. Today's drill: Answer 3 hard questions with zero qualifiers.'
- **Quick Stats Row:** Sessions this week | Current streak | Top skill
- **Upcoming Moment:** Countdown to your real event. Creates urgency.

> *No clutter. No 'explore' or 'discover.' You open the app to practice or be coached.*

3.2 Practice — The Sparring Room Prep

Feel: Like suiting up. Dark, focused, slightly tense.

Step 0 — Choose Mode:

- **With Material** — Upload deck, doc, or paste text
- **Without Material** — Speak to a topic from memory

Step 1 — What are you defending? Drop zone for files (PPTX, PDF, TXT). AI auto-extracts 3-5 key claims.

Step 2 — What's the room? Pick a scenario: Thesis Defense, Investor Pitch, Board Update, Performance Review, Grant Defense, Client Presentation, Custom.

Step 3 — Configure Your Panel. 2-3 seats. Each = an archetype:

- **The Skeptic** — questions your core assumption
- **The Nitpicker** — dives into methodology and details
- **The So-What** — challenges your impact and relevance
- **The Timekeeper** — cuts you off, forces concision
- **The Ally** — seems supportive but probes for weakness

Each has: editable name, behavior style (Gentle/Moderate/Aggressive), voice preview.

Step 4 — Start Sparring. Big button. Full screen transition. You're in the room.

3.3 The Sparring Room (Live Session)

Feel: You're in a dark room, presenting. The panel is watching.

- **Left 60%:** Your content (slide viewer or text outline). Auto-advances or manual.
- **Right 40%:** The Panel — 2-3 persona cards. Each shows: Listening | Speaking | Thinking. Live mood indicator. Who's about to interrupt — subtle pulse animation.
- **Bottom Bar:** Slide controls, Mute/Pause, **'I need a moment'** (pauses panel, coach whispers a framework), Timer, End Session
- **Live Telemetry (subtle):** WPM, filler count, pause time — numbers that float in when relevant, not a dashboard

The Magic: The panel interrupts mid-sentence. They reference your actual slides. They remember what you said 3 minutes ago and call out contradictions. They get harder if you're breezing through.

3.4 Post-Session Debrief (Where Coaching Lives)

Feel: Like sitting down with your coach after a spar. Honest, specific, forward-looking.

Section 1: The Moment — 'You were challenged on your revenue projection at 4:32. You paused for 3 seconds, then reframed with the cohort data. That was your best recovery this week.'

Section 2: What You Did Well — 2-3 specific, evidence-based observations.

Section 3: The Pattern — 'This is the third session where you over-explain when questioned on methodology. Try: Acknowledge the concern in one sentence, then answer.'

Section 4: One Thing to Practice — A single, actionable focus. Not a list of 10.

Section 5: Telemetry (Framed as Insight) — 'Pace: 142 WPM (rushing first 5 min, settled after). Fillers: 4/min (down from 6). Recovery: 3.2s avg (target: <2s).'

Section 6: Save or Share — Export summary to send to a human mentor/coach.

3.5 Coach Screen (The Development Layer)

Feel: Like a training curriculum. Structured, intentional.

Tab 1: Programs — Active sprints. Example: *Thesis Defense Ready* (Week 2 of 4), *Pitch Confidence* (Not started). Each has 4 structured sessions with clear before/after.

Tab 2: Skill Map — Radar chart or bars across 5 dimensions:

1. **Clarity** — Do you get to the point?
2. **Recovery** — How do you handle interruption?
3. **Concision** — Do you respect the room's time?
4. **Presence** — Vocal confidence, pace, pause
5. **Evidence** — Do you ground claims in proof?

Tab 3: Insights — Running list of pattern notes. 'You tend to hedge with 'I think' when challenged. Work on declarative framing.'

3.6 Progress Screen

Feel: Like a performance review on yourself. Data-rich but human.

- **Communication Profile:** AI-generated paragraph summarizing you as a communicator.
- **Session History:** List view with key insight from each session.
- **Growth Trends:** Line charts showing filler word reduction, recovery time improvement.
- **Milestones:** Tasteful. Not badges. '10 sessions completed — you've practiced more than 80% of users.'

4. THE COACHING LOOP (Baked In, Not Bolted On)

This is the architecture that makes it a coach, not a chatbot:

Every Session → AI generates: What happened + What was hard + What you did well

↓ Stores in your Pattern Memory

↓ Next session: AI reads last 3 summaries before starting

↓ Panel behavior adjusts (e.g., 'They struggle with methodology, probe there')

↓ Post-session: Compare to previous. Show movement.

↓ Coach tab: Update Skill Map + add Insight + recommend next focus

The result: After 3 sessions, the user feels like SharkPit *knows them.* That's the premium feeling.

5. THE 'NOT TOYISH' FIXES

What's making it feel toyish	The fix
Anime/cartoon avatars	Abstract geometric portraits or clean initials with color coding. Professional, not playful.
'Professor Aris' naming	Default to archetypes: 'The Skeptic.' Let users rename if they want, but don't force whimsy.
Gamey language ('level up', 'xp')	Remove all of it. Use coaching language: 'Sprint,' 'Focus,' 'Pattern,' 'Baseline.'
Bright, playful colors	Stick to deep charcoal, slate, and one accent blue. No neon. No gradients for fun.
Metrics as a game HUD	Telemetry should feel like a coach's notes — contextual, sparse, insightful.
'Rehearse' language	Use 'Practice' or 'Spar.' 'Rehearse' implies theater. 'Practice' implies sport.

6. BUILD ORDER

One concrete thing at a time. No parallel tracks.

Phase 1: The Sparring Engine (You mostly have this)

- Upload deck/topic
- Configure 3 archetype personas
- Live voice session with interruption
- Basic telemetry capture

Phase 2: The Coaching Layer (The Differentiator)

- Post-session AI insight generation (the 4-section debrief)
- Pattern memory (store summaries, feed into next session prompt)
- Skill tracking (5 dimensions, simple scoring)
- 'Today's Focus' on Home screen

Phase 3: Programs (The Structured Path)

- Build ONE program: 'Thesis Defense Sprint' (4 sessions, structured)
- Hardcode it beautifully. Don't build a program builder yet.
- Active program card on Home + Coach tab

Phase 4: Polish & Premium Feel

- Redesign personas (remove cartoon vibes)
- Voice quality upgrade (if needed)
- Advanced telemetry with interpretation
- Share summary feature

7. THE 1-WEEK SPRINT — ZERO AMBIGUITY

Goal: Ship v1.0 within 7 days. **Rule:** Every task has an owner, a deliverable, and a deadline.

7.1 Roles

YOU (Engineering)	HIM (Methodology)
Builds the product. Ships code. Owns infra, UI, AI pipeline, telemetry.	Owns the soul of the product. Writes what the AI says, when it interrupts, how it coaches.

The Handoff Rule: He writes documents. You turn documents into code. No exceptions.

DAY 1 — MONDAY: Failure Taxonomy

His task (4h): Find 3 real breakdown videos. Timestamp the collapse. Write 1 paragraph per breakdown: trigger, failure, fix.

Your task (4h): Set up shared Notion workspace. Build crude 'Methodology Admin Panel' for him to edit persona prompts.

Deliverable by 11:59 PM: Shared doc with 5-7 failure categories named and described.

DAY 2 — TUESDAY: Persona Archetypes

His task (4h): Write full character sheets for The Skeptic, The Nitpicker, The Ally. Include triggers, question templates, aggression levels, forbidden behaviors. Record 2-min voice samples for each.

Your task (4h): Integrate voice samples into TTS pipeline. Wire persona prompts into live session system prompt.

Deliverable by 11:59 PM: 3 working personas with distinct voices and behaviors.

DAY 3 — WEDNESDAY: The Debrief Rubric

His task (4h): Build 5-dimension scoring rubric in Google Sheets. For each dimension: 3 specific markers, good/bad examples, severity, coach tip.

Your task (4h): Build post-session debrief screen UI. Implement basic telemetry capture (WPM, filler count, pause time, recovery time).

Deliverable by 11:59 PM: Working rubric spreadsheet + debrief screen that displays telemetry.

DAY 4 — THURSDAY: 'I Need a Moment' Frameworks

His task (3h): Write 5 intervention scripts for common failure modes. Each: trigger condition, coach whisper text, example scenario.

Your task (5h): Build the 'I Need a Moment' button in live session. Build classifier that matches failure patterns to right intervention script.

Deliverable by 11:59 PM: Working panic button with 5 real coach interventions.

DAY 5 — FRIDAY: First Program Curriculum

His task (4h): Design ONE 4-session program: 'High-Stakes Defense Sprint.' Each session: objective, persona config, time limit, expected outcome. Write before/after comparison narrative.

Your task (4h): Build Coach tab with Programs section. Build session scheduling and progress tracking.

Deliverable by 11:59 PM: One hardcoded program that users can enroll in and progress through.

DAY 6 — SATURDAY: Integration & Polish

Both of you (4h each): He tests full flow: upload → configure → spar → debrief → progress. Records Loom videos for every bug or 'feels fake' moment. You fix critical bugs only. No new features.

Deliverable by 11:59 PM: Zero critical bugs. Product is end-to-end usable.

DAY 7 — SUNDAY: Ship Day

His task (3h): Write landing page copy (headline, subhead, 3 value props, CTA). Record 90-second demo video showing full flow.

Your task (5h): Deploy to production. Set up analytics (session count, completion rate, NPS). Send product to 5 beta users for feedback.

Deliverable by 11:59 PM: Live URL. Demo video. 5 beta invites sent.

8. EXACT DELIVERABLE TEMPLATES

8.1 Persona Character Sheet Template

Persona Name: [e.g., The Skeptic]
Archetype: [Skeptic / Nitpicker / Ally / So-What / Timekeeper]
Voice ID: [Link to voice sample file]
Speaking pace: [Slow / Normal / Fast]
Tone: [Warm / Neutral / Cold / Confrontational]

Trigger Conditions (when to interrupt):
1. [Condition 1]
2. [Condition 2]
3. [Condition 3]

Question Templates:
1. "What evidence supports [claim]?"
2. "Have you considered [alternative]?"
3. "That assumes [X] — what if [X] is wrong?"

Aggression Levels:
• Gentle: "[soft version]"
• Moderate: "[direct version]"
• Aggressive: "[sharp version]"

Forbidden Behaviors:
1. Never attack the person. Only the argument.
2. Never ask two questions at once.
3. If user recovers well, acknowledge briefly then move on.

Example Dialogue:
User: "Our market is huge."
Skeptic (Moderate): "Huge is vague. What's your TAM, and what percentage can you realistically capture?"

8.2 Debrief Rubric Template

Dimension: [Clarity / Recovery / Concision / Presence / Evidence]
Marker: [Exact behavior to detect]
Detection method: [Keyword / Timing / NLP / Manual]
Good example: [What success looks like]
Bad example: [What failure looks like]
Severity: [High / Medium / Low]
Coach tip: [What to tell the user]

Example:
Dimension: Hedging (under Presence)
Marker: Uses "I think," "maybe," "sort of" when challenged
Detection: Keyword regex
Good: "Revenue will grow 40%"

Bad: "I think revenue might grow maybe like 40%"
Severity: High if >3 per minute
Coach tip: "Drop the qualifiers. Say it directly. You can be wrong, but be clear."

8.3 Intervention Script Template

Intervention Name: [e.g., ACKNOWLEDGE → REFRAME → ANSWER]
Trigger: [When does the coach whisper this?]
Script:
1. "Pause. Take a breath." [pause 2s]
2. "Acknowledge the concern: 'That's a fair point.'"
3. "Reframe the scope: 'What you're raising is exactly why we focused on...'"
4. "Answer with one piece of evidence. Then stop."

Example scenario:
User is challenged: "Your sample size seems small."
Coach whispers: "Pause. Say: 'You're right to flag that. We addressed it by running a power analysis.' Then give the number. One sentence. Stop."

8.4 QA Report Template

Issue #: [Auto-increment]
Date: [Date]
Video: [Loom link]
Timestamp: [MM:SS]
Persona: [Which archetype]
Problem: [What happened]
Expected: [What should have happened]
Severity: [Blocker / High / Medium / Low]
Your fix: [You fill this in after fixing]
Status: [Open / In Progress / Fixed / Won't Fix]

9. THE DEBATE COACH ANGLE

Can SharkPit work as a debate coach? Will debaters find it useful? The honest answer: **Yes, but not yet as a pure debate coach. Here's the roadmap.**

9.1 What SharkPit Already Does for Debaters

Debate is fundamentally **defense under pressure** — exactly what SharkPit trains.

- **Cross-examination practice:** The Skeptic persona behaves like a CX opponent — probing assumptions, asking for evidence, catching contradictions.
- **Rebuttal pressure-testing:** Upload your case brief. The AI attacks your weakest contention. You practice defending it live.
- **Recovery training:** Debaters lose rounds not because their case is weak, but because they freeze when caught off-guard. The interruption mechanic trains exactly this.
- **Concision under fire:** The Timekeeper persona forces 15-second answers. This is pure rebuttal prep.
- **Pattern recognition:** After 5 sessions, the AI knows you over-explain on impact calculus. That's coaching gold.

9.2 What Debate-Specific Features Would Be Needed (Post-v1)

Feature	What It Does	Priority
Flow integration	AI generates a flow of your speech and attacks the drops	High
Argument type detection	AI recognizes DA, CP, K, T, case args and responds appropriately	High
Speed/Spread mode	AI listens at 300+ WPM and asks about specific cards you read	High
Evidence challenge	AI asks 'What's your warrant?' or 'Date?' or 'Author quals?'	Medium
Judge adaptation	Personas simulate lay, flow, and kritikal judges	Medium
Partner practice	Two debaters practice together with AI as judge/opponent	Low

9.3 The Market for Debate Coaching

Current state: Private debate coaches charge $75-$200/hour. Top-tier programs cost $3,000-$5,000 per summer. Most debaters cannot afford this.

The gap: A debater needs 20-50 hours of practice rounds per topic. At $100/hr, that's $2,000-$5,000 per semester. SharkPit at $29/month replaces 80% of that practice for 1% of the cost.

The urgency: Tournament season is predictable. You can market seasonally: 'Topic release is Monday. Start drilling with SharkPit by Wednesday.'

9.4 The Verdict

For v1 (this week): SharkPit is a general high-stakes communication coach that happens to be excellent for debaters. Your friend should use it, test it, and give feedback. But don't market it as 'for debate' yet.

For v2 (Month 2-3): Add flow integration and argument-type detection. Then launch a 'Debate Sparring' mode. Same engine, debate-specific landing page.

For v3 (Month 4-6): Partner with a debate circuit or camp. Offer team licenses. Become the standard practice tool for programs that can't afford a full-time coach.

> *Bottom line: Debaters are your perfect early adopters. They understand the value of sparring immediately. They practice obsessively. They talk to each other. But don't narrow your product to debate until the core engine is bulletproof.*

10. CO-FOUNDER ROLES & RESPONSIBILITIES

10.1 His Role: Chief Methodology Officer

He owns everything that happens **inside the sparring room.** You own everything that happens **outside and around it.**

YOU (Engineering)	HIM (Methodology)
The voice AI pipeline	What the AI actually says and when
The UI, sidebar, telemetry	The 5 skill dimensions and how they're scored
Upload, parsing, slide sync	The persona archetypes (Skeptic, Nitpicker, etc.)
Session storage, pattern memory	The pattern recognition logic — what counts as 'hedging'
The 'I need a moment' button	What the coach whispers when the button is pressed
Payment, auth, infra	The curriculum — Week 1 to Week 4 progression
Demo videos, landing page code	The landing page story — how you sell the pain

10.2 What He Actually Does (No Code Required)

1. Persona Design — Google Docs / Notion

He writes character sheets. You turn them into system prompts.

2. The Debrief Rubric — Google Sheets

He defines what 'good' looks like across 5 dimensions. You build the detection logic.

3. Intervention Scripts — Text Files

He writes what the coach whispers. You build the classifier that triggers the right one.

4. Curriculum Design — Notion / Linear

He designs program structure. You build it into the product.

5. Reference Audio — Voice Recorder

He records voice samples for each archetype. You feed them to TTS.

6. QA Testing — Loom Videos + GitHub Issues

He uses the product. Records 2-min Loom videos for every bug. You fix them.

7. User Research — Interviews + Shared Notes

He conducts user interviews. You both read them and look for recurring themes.

10.3 The Founders' Agreement (Have This Conversation This Week)

Equity: 30-40% to him. Vesting over 4 years, 1-year cliff.
Time commitment: Define minimum hours per week (e.g., 15 hours for part-time).

Decision rights: You own technical decisions. He owns methodology decisions. Everything else: decide together. Deadlock = coin flip or advisor.
IP: Everything built belongs to the company.
Exit clause: If someone leaves after 6 months, the other has right to buy back equity at pre-agreed formula.

Write it down. Sign it. 30 minutes. Saves 3 years of resentment.

11. EXECUTIVE COACHING MARKET RESEARCH

11.1 Market Size

Source	2026 Market Size	Notes
Mordor Intelligence	$112.98 billion	Includes full leadership development + coaching programs
ICF Global Coaching Study	$5.34 billion	Practitioner revenue only (narrower scope)
Executive Coaching Certification	$14.1 billion	Certification + training market
Projected 2031	$174.53 billion	9.11% CAGR

11.2 Proven ROI

- **Average ROI: 5-7x** the investment — for every $1 spent, companies make $5-$7 back
- **86%** of companies that calculated ROI made back their initial investment
- **70%** of Fortune 500 companies use executive coaching
- **99%** of clients are 'satisfied' or 'very satisfied'
- **70%** of individuals report improved work performance after coaching
- **51%** of companies using coaching report higher revenue than other businesses

11.3 The Urgency Is Real

- **81%** of C-suite executives rate improving mental health above advancing their career
- **70%** of C-suite executives report burnout and are seriously considering leaving
- C-suite succession gaps are driving sustained budget growth
- Corporate strategy now turns over faster than traditional development programs can run
- When senior roles are hard to fill, developing who you have stops being optional

11.4 The Supply-Demand Mismatch (Your Opportunity)

- **54%** of coaches say better platforms are a priority — but only **19%** invested in new technology last year
- Only **47%** of coaches use a digital platform at all
- Most coaching is still **in-person** (56.62% of market)
- Average coaching rate: **$234/hour** — but average active coach earns only **$49,283/year**
- **SMEs** growing at **11.17% CAGR** — precisely the customers who cannot afford $500/hr human coaches

11.5 The Bottom Line

Executive coaching is **not a luxury anymore.** It's workforce infrastructure. The individual market — mid-level managers, early-career founders, grad students entering the workforce — is the biggest whitespace. They need coaching. They can't pay $234/hour. They'll pay $29-79/month for software that delivers 60% of the value.

12. ENTERTAINMENT DEMOCRATIZATION IDEAS

Ultra-premium entertainment experiences that software can make accessible to everyone.

Idea	Ultra-Rich Version	Democratized Version	Software Lever
Private Concerts	$100k-$1M for celebrity private gig	$5-20 for intimate livestream shows (50-500 people)	Livestream platform + artist marketplace
Red Carpet Home Cinema	$15k server + $2-3k per movie	$30-50 day-and-date streaming rental	Premium streaming layer
VIP Access Marketplace	Invitation-only backstage, meet-and-greets	$10-15 for virtual backstage passes, creator subs	Creator-fan proximity platform
Celebrity Micro-Experiences	$100k-$1M private appearances	$15-50 for 5-min live calls, 'cook with me' sessions	Scheduling + video infrastructure
Paid Digital Clubs	Members-only clubs, private islands	$20-50/mo for curated, ad-free communities	Community platform + event tooling

Honest take: Entertainment is fun but crowded. Coaching is where you have product, a co-founder with domain expertise, and a $113B market with proven ROI. Build SharkPit first. Entertainment can be Phase 2.

13. SHARED WORKSPACE SETUP

Set this up on Day 1. No exceptions.

Notion Workspace Structure:

- **01 — Research** (Failure taxonomy, user interview notes, competitive analysis)
- **02 — Methodology** (Persona docs, rubric sheets, intervention scripts, curriculum)
- **03 — QA** (Bug reports, Loom videos, feature requests)
- **04 — Product** (Screenshots, demo videos, landing page copy)
- **05 — Metrics** (Weekly stats, user feedback, NPS scores)

Communication Rules:

- All persona changes → go in the Notion doc first. You pull from the doc.
- All bugs → Loom video + Notion QA database entry. No Slack DMs about bugs.
- Daily standup: 15 min voice call. What did you ship yesterday? What are you shipping today? What's blocked?
- No scope creep during the week. If he has a new idea, it goes in 'Week 2 Ideas' and stays there.

His Weekly Rhythm (Once Building)

Monday (1h): Reviews weekend QA reports, updates persona docs

Tuesday (2h): Conducts 1 user interview, writes notes

Wednesday (1h): Tests latest build, records Loom bugs

Thursday (2h): Writes new intervention scripts or curriculum content

Friday (1h): Weekly sync with you — reviews metrics, plans next week

Total: ~7 hours/week on the side. If full-time: 40+ hours of deep methodology work, user research, and sales storytelling.